

## Task 7.1

### 7.1: Elasticity Task

#### Integrating LVM with Hadoop and providing Elasticity to DataNode Storage

```
https://aws.amazon.com/amazon-linux-2/
25 package(s) needed for security, out of 39 available
Run "sudo yum update" to apply all updates.
[ec2-user@ip-172-31-7-110 ~]$ sudo su - root
Last login: Fri Oct 30 07:24:37 UTC 2020 on pts/0
[root@ip-172-31-7-110 ~]# fdisk -l
Disk /dev/xvda: 8 GiB, 8589934592 bytes, 16777216 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: 66B3909F-969E-4FD1-901C-CEE3A9974A83
```

Device	Start	End	Sectors	Size	Type
/dev/xvda1	4096	16777182	16773087	8G	Linux filesystem
/dev/xvda128	2048	4095	2048	1M	BIOS boot

Partition table entries are not in disk order.

```
Disk /dev/xvdf: 20 GiB, 21474836480 bytes, 41943040 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
```

```
Disk /dev/xvdg: 30 GiB, 32212254720 bytes, 62914560 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
[root@ip-172-31-7-110 ~]#
```

Attached two separate  
EBS storage(Hard disks)  
of size 20 GB and 30 GB  
with device names as  
/dev/xvdf - HD1  
/dev/xvdg - HD2

```
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
```

```
[root@ip-172-31-7-110 ~]# pvcreate /dev/xvdf
Physical volume "/dev/xvdf" successfully created.
```

```
[root@ip-172-31-7-110 ~]# pvcreate /dev/xvdg
Physical volume "/dev/xvdg" successfully created.
```

```
[root@ip-172-31-7-110 ~]# pvdisplay
```

```
"/dev/sdf" is a new physical volume of "20.00 GiB"
```

```
--- NEW Physical volume ---
```

PV Name	/dev/sdf
VG Name	
PV Size	20.00 GiB
Allocatable	NO
PE Size	0
Total PE	0
Free PE	0
Allocated PE	0
PV UUID	6tx4cx-FWNw-5FxFZ-zpJi-X7EI-C4t7-2HWff9

Not allocated to any  
VG yet.

```
"/dev/sdg" is a new physical volume of "30.00 GiB"
```

```
--- NEW Physical volume ---
```

PV Name	/dev/sdg
VG Name	
PV Size	30.00 GiB
Allocatable	NO
PE Size	0
Total PE	0
Free PE	0
Allocated PE	0
PV UUID	UL0CFM-IdXv-omXd-uNqV-UjtX-75mj-znDCTN

```
[root@ip-172-31-7-110 ~]# D
```

```
[root@ip-172-31-7-110 ~]# pvdisplay /dev/sdf ➡ Showing info of PV1
```

```
"/dev/sdf" is a new physical volume of "20.00 GiB"
```

```
--- NEW Physical volume ---
```

PV Name	/dev/sdf
VG Name	
PV Size	20.00 GiB
Allocatable	NO
PE Size	0
Total PE	0
Free PE	0
Allocated PE	0
PV UUID	6tx4cx-FWNw-5FxZ-zpJi-X7EI-C4t7-2HWff9

```
[root@ip-172-31-7-110 ~]# pvdisplay /dev/sdg ➡ Showing info of PV2
```

```
"/dev/sdg" is a new physical volume of "30.00 GiB"
```

```
--- NEW Physical volume ---
```

PV Name	/dev/sdg
VG Name	
PV Size	30.00 GiB
Allocatable	NO
PE Size	0
Total PE	0
Free PE	0
Allocated PE	0
PV UUID	UL0CFM-IdXv-omXd-uNqV-UjtX-75mj-znDCTN

```
[root@ip-172-31-7-110 ~]#
```

```
[root@ip-172-31-7-110 ~]# vgdisplay ➡ Showing all available VGs
```

```
[root@ip-172-31-7-110 ~]# vgdisplay iiecvg ➡ Checking if any VG with this  
Volume group "iiecvg" not found name 'iiecvg' available
```

```
Cannot process volume group iiecvg
```

```
[root@ip-172-31-7-110 ~]# vgcreate iiecvg /dev/sdf /dev/sdg ➡ Creating a VG with  
Volume group "iiecvg" successfully created combined storage of  
PV1 and PV2
```

```
[root@ip-172-31-7-110 ~]# vgdisplay iiecvg ➡ Info related to VG
```

```
--- Volume group ---
```

VG Name	iiecvg
System ID	
Format	lvm2
Metadata Areas	2
Metadata Sequence No	1
VG Access	read/write
VG Status	resizable
MAX LV	0
Cur LV	0
Open LV	0
Max PV	0
Cur PV	2
Act PV	2
VG Size	49.99 GiB
PE Size	4.00 MiB
Total PE	12798
Alloc PE / Size	0 / 0
Free PE / Size	12798 / 49.99 GiB
VG UUID	uhk8GW-sacZ-VTyw-BrjU-LtrT-rMSP-CdSjGp

➡ Space available for storing is not allocated yet. First create partition to store.

➡ Total storage 20GB + 30GB nearly 50GB is unallocated

```
[root@ip-172-31-7-110 ~]#
```

```
[root@ip-172-31-7-110 ~]# lvcreate --size 40G --name lv1 iiecvg
Logical volume "lv1" created.
[root@ip-172-31-7-110 ~]# vgdisplay iiecvg
--- Volume group ---
VG Name                iiecvg
System ID
Format                 lvm2
Metadata Areas         2
Metadata Sequence No   2
VG Access               read/write
VG Status               resizable
MAX LV                 0
Cur LV                1
Open LV                0
Max PV                 0
Cur PV                2
Act PV                2
VG Size                49.99 GiB
PE Size                4.00 MiB
Total PE              12798
Alloc PE / Size        10240 / 40.00 GiB
Free PE / Size         2558 / 9.99 GiB
VG UUID                uhk8GW-sacZ-VTyw-BrjU-LtrT-rMSP-CdSjGp
```

⇒ Creating a dynamic partition called LV(Logival Volume) inside VG of size 40GB getting storage from both PVs.

⇒ Total allocated size i.e. space available to store is 40 GB(partition size)

⇒ Unallocated Space in VG

```
[root@ip-172-31-7-110 ~]#
```

```
[root@ip-172-31-7-110 ~]# lvdisplay iiecvg/lv1
--- Logical volume ---
LV Path                /dev/iiecvg/lv1
LV Name                lv1
VG Name                iiecvg
LV UUID                VITKJc-QSuA-0Q2W-YJP2-zPKI-uq2e-UDKA9F
LV Write Access        read/write
LV Creation host, time ip-172-31-7-110.us-east-2.compute.internal, 2020-10-30 07:56:11 +0000
LV Status              available
# open                 0
LV Size                40.00 GiB
Current LE             10240
Segments               2
Allocation              inherit
Read ahead sectors     auto
- currently set to    256
Block device           253:0
```

⇒ Info related to LV

```
[root@ip-172-31-7-110 ~]# lvdisplay
--- Logical volume ---
LV Path                /dev/iiecvg/lv1
LV Name                lv1
VG Name                iiecvg
LV UUID                VITKJc-QSuA-0Q2W-YJP2-zPKI-uq2e-UDKA9F
LV Write Access        read/write
LV Creation host, time ip-172-31-7-110.us-east-2.compute.internal, 2020-10-30 07:56:11 +0000
LV Status              available
# open                 0
LV Size                40.00 GiB
Current LE             10240
Segments               2
Allocation              inherit
Read ahead sectors     auto
- currently set to    256
Block device           253:0
```

```
[root@ip-172-31-7-110 ~]# mkfs.ext4 /dev/iiecvq/lv1 ➡ Formatting the LV
mke2fs 1.42.9 (28-Dec-2013)
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=0 blocks, Stripe width=0 blocks
2621440 inodes, 10485760 blocks
524288 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=2157969408
320 block groups
32768 blocks per group, 32768 fragments per group
8192 inodes per group
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632, 2654208,
    4096000, 7962624

Allocating group tables: done
Writing inode tables: done
Creating journal (32768 blocks): done
Writing superblocks and filesystem accounting information: done
```

```
[root@ip-172-31-7-110 hadoop]# hadoop dfsadmin -report
Safe mode is ON
Configured Capacity: 8577331200 (7.99 GB)
Present Capacity: 6481682432 (6.04 GB)
DFS Remaining: 6481674240 (6.04 GB)
DFS Used: 8192 (8 KB)
DFS Used%: 0%
Under replicated blocks: 0
Blocks with corrupt replicas: 0
Missing blocks: 0

-----
Datanodes available: 1 (1 total, 0 dead)

Name: 18.223.239.112:50010
Decommission Status : Normal
Configured Capacity: 8577331200 (7.99 GB) ➡ Initially the data node
DFS Used: 8192 (8 KB) was sharing the storage
Non DFS Used: 2095648768 (1.95 GB) of its Root or Slash
DFS Remaining: 6481674240 (6.04 GB) drive via /dn directory
DFS Used%: 0% nearly equal to 8 GB.
DFS Remaining%: 75.57%
Last contact: Fri Oct 30 08:20:57 UTC 2020

[root@ip-172-31-7-110 hadoop]#
```

```
[root@ip-172-31-7-110 hadoop]# mount /dev/iiecvlg/lv1 /dn
[root@ip-172-31-7-110 hadoop]# df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
devtmpfs	474M	0	474M	0%	/dev
tmpfs	492M	0	492M	0%	/dev/shm
tmpfs	492M	508K	492M	1%	/run
tmpfs	492M	0	492M	0%	/sys/fs/cgroup
/dev/xvda1	8.0G	2.0G	6.1G	25%	/
tmpfs	99M	0	99M	0%	/run/user/1000
tmpfs	99M	0	99M	0%	/run/user/1000
/dev/mapper/iiecvlg-lv1	40G	49M	38G	1%	/dn

```
[root@ip-172-31-7-110 hadoop]#
```

Mounting the created LV to the /dn directory shared by data node to Master node.

After mount the size of /dn becomes near to 40GB as now getting space from LV instead of Root device.

```
[root@ip-172-31-7-110 hadoop]# hadoop dfsadmin-report
Error: Could not find or load main class dfsadmin-report
[root@ip-172-31-7-110 hadoop]# hadoop dfsadmin -report
Safe mode is ON
Configured Capacity: 42141450240 (39.25 GB)
Present Capacity: 39926841344 (37.18 GB)
DFS Remaining: 39926833152 (37.18 GB)
DFS Used: 8192 (8 KB)
DFS Used%: 0%
Under replicated blocks: 0
Blocks with corrupt replicas: 0
Missing blocks: 0

-----
Datanodes available: 1 (1 total, 0 dead)

Name: 18.223.239.112:50010
Decommission Status : Normal
Configured Capacity: 42141450240 (39.25 GB)
DFS Used: 8192 (8 KB)
Non DFS Used: 2214608896 (2.06 GB)
DFS Remaining: 39926833152 (37.18 GB)
DFS Used%: 0%
DFS Remaining%: 94.74%
Last contact: Fri Oct 30 08:23:45 UTC 2020

[root@ip-172-31-7-110 hadoop]#
```

Total storage size shared to master node is 40 GB



```
[root@ip-172-31-7-110 hadoop]# lvdisplay
```

```

--- Logical volume ---
LV Path                /dev/iiecvg/lv1
LV Name                 lv1
VG Name                 iiecvg
LV UUID                 VITKJc-QSuA-0Q2W-YJP2-zPKI-uq2e-UDKA9F
LV Write Access         read/write
LV Creation host, time ip-172-31-7-110.us-east-2.compute.internal, 2020-10-30 07:56:11 +0000
LV Status                available
# open                  1
LV Size                40.00 GiB
Current LE              10240
Segments                2
Allocation               inherit
Read ahead sectors      auto
- currently set to      256
Block device            253:0

```

#df -h command shows the total available size capable to store data(formatted space). Whereas #lvdisplay command shows the total size.

```
[root@ip-172-31-7-110 hadoop]# df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
devtmpfs	474M	0	474M	0%	/dev
tmpfs	492M	0	492M	0%	/dev/shm
tmpfs	492M	508K	492M	1%	/run
tmpfs	492M	0	492M	0%	/sys/fs/cgroup
/dev/xvda1	8.0G	2.0G	6.1G	25%	/
tmpfs	99M	0	99M	0%	/run/user/1000
tmpfs	99M	0	99M	0%	/run/user/0
/dev/mapper/iiecvg-lv1	40G	49M	<b>38G</b>	1%	/dn

```
[root@ip-172-31-7-110 hadoop]#
```

```
[root@ip-172-31-7-110 hadoop]# lvextend --size +5G /dev/iiecvg/lv1
```

```
Size of logical volume iiecvg/lv1 changed from 40.00 GiB (10240 extents) to 45.00 GiB (11520 extents).
```

```
Logical volume iiecvg/lv1 successfully resized.
```

```
[root@ip-172-31-7-110 hadoop]# lvdisplay /dev/iiecvg/lv1
```

```

--- Logical volume ---
LV Path                /dev/iiecvg/lv1
LV Name                 lv1
VG Name                 iiecvg
LV UUID                 VITKJc-QSuA-0Q2W-YJP2-zPKI-uq2e-UDKA9F
LV Write Access         read/write
LV Creation host, time ip-172-31-7-110.us-east-2.compute.internal, 2020-10-30 07:56:11 +0000
LV Status                available
# open                  1
LV Size                45.00 GiB
Current LE              11520
Segments                2
Allocation               inherit
Read ahead sectors      auto
- currently set to      256
Block device            253:0

```

Here showing the total size of LV as 45GB after extension.

Demonstrating the elasticity feature of the LV shared via /dn directory to Hadoop master. We can extend the size of the LV.

```
[root@ip-172-31-7-110 hadoop]# df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
devtmpfs	474M	0	474M	0%	/dev
tmpfs	492M	0	492M	0%	/dev/shm
tmpfs	492M	508K	492M	1%	/run
tmpfs	492M	0	492M	0%	/sys/fs/cgroup
/dev/xvda1	8.0G	2.0G	6.1G	25%	/
tmpfs	99M	0	99M	0%	/run/user/1000
tmpfs	99M	0	99M	0%	/run/user/0
/dev/mapper/iiecvg-lv1	40G	49M	<b>38G</b>	1%	/dn

```
[root@ip-172-31-7-110 hadoop]#
```

But here the extension is not reflected and available size is still nearly 40GB. This because extended part of LV is not formatted, so incapable of storing data.

```
[root@ip-172-31-7-110 hadoop]# resize2fs /dev/iiecvlg/lv1
resize2fs 1.42.9 (28-Dec-2013)
Filesystem at /dev/iiecvlg/lv1 is mounted on /dn; on-line resiz
old_desc_blocks = 5, new_desc_blocks = 6
The filesystem on /dev/iiecvlg/lv1 is now 11796480 blocks long
```

⇒ Foramtting only the extended part without data removal from previously formatted part of LV.

```
[root@ip-172-31-7-110 hadoop]# df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
devtmpfs	474M	0	474M	0%	/dev
tmpfs	492M	0	492M	0%	/dev/shm
tmpfs	492M	508K	492M	1%	/run
tmpfs	492M	0	492M	0%	/sys/fs/cgroup
/dev/xvda1	8.0G	2.0G	6.1G	25%	/
tmpfs	99M	0	99M	0%	/run/user/1000
tmpfs	99M	0	99M	0%	/run/user/0

```
/dev/mapper/iiecvlg-lv1 45G 52M 42G 1% /dn
```

Now available size extended after formatting.

```
[root@ip-172-31-7-110 hadoop]#
```

```
[root@ip-172-31-7-110 hadoop]# hadoop dfsadmin -report
```

```
Safe mode is ON
Configured Capacity: 47425867776 (44.17 GB)
Present Capacity: 44970913792 (41.88 GB)
DFS Remaining: 44970893312 (41.88 GB)
DFS Used: 20480 (20 KB)
DFS Used%: 0%
Under replicated blocks: 0
Blocks with corrupt replicas: 0
Missing blocks: 0
```

```
-----
Datanodes available: 1 (1 total, 0 dead)
```

```
Name: 18.223.239.112:50010
Decommission Status : Normal
Configured Capacity: 47425867776 (44.17 GB)
DFS Used: 20480 (20 KB)
Non DFS Used: 2454953984 (2.29 GB)
DFS Remaining: 44970893312(41.88 GB)
DFS Used%: 0%
DFS Remaining%: 94.82%
Last contact: Fri Oct 30 08:34:45 UTC 2020
```

⇒ After extending the size of LV, the shared storage also extended achieving the functionality of elasticity. So on integrating LVM with Hadoop, provided Elasticity to DataNode Storage

```
[root@ip-172-31-7-110 hadoop]#
```